

BUILDING THE MOST ADVANCED RETRIEVAL TOOL FOR AI AGENTS TO ENHANCE RELIABILITY AND CONTEXTUAL INTELLIGENCE THROUGH AN AI-POWERED RAG API.


• Developer & IT Infrastructure > AI Developer Infrastructure
• B2D > SaaS

PRE-SCREENING SCORE
Thesis : Profund VC

NOTE: This is a raw pre-screening score. Thesis weights are applied in the Synthetic GP qualification pipeline after angle detection.

TEAM EXCELLENCE 5/100
MARKET OPPORTUNITY 85/100
PRODUCT INNOVATION 80/100
BUSINESS MODEL 75/100
TRACTION & GROWTH 60/100

PRE-SCREENING SCORE: 61/100 → WEAK SIGNAL



? In a NUTSHELL : Laive.ai is a AI Developer Infrastructure that enables AI Developers to build reliable, context-aware LLM applications by using a managed, high-performance RAG API.

! The PROBLEM : An AI developer building a chatbot needs to connect it to their company's internal documents, but stitching together open-source vector databases, embedding models, and chunking logic results in a fragile, slow, and unreliable system that frequently gives incorrect answers.

✓ The SOLUTION : Laive.ai provides a single, fully managed API endpoint that ingests raw enterprise data from multiple sources and uses a sophisticated multi-layer retrieval pipeline to instantly find the precise information an LLM needs to answer a query accurately.

🚀 The GTM : The company is pursuing a product-led growth motion targeting individual AI developers, offering a freemium tier as the hook. This is the smartest entry point as it seeds the market and establishes a foothold within engineering teams, creating internal champions who can later drive larger, more lucrative enterprise-wide adoption.

👤 TEAM EXCELLENCE (0%) | Score: 5/100
This score is severely suppressed by a complete lack of public data; it does not necessarily reflect poor quality but rather total opacity, which is a significant risk.

- Founder-Market Fit (25%) | Score: 0/100: No data is available to evaluate the CEO, Florian Palmade; there is no information on his industry experience or 'Earned Secret'.
- Track Record (25%) | Score: 0/100: No previous exits, awards, or notable investor backing are documented in the provided data.
- Leadership (25%) | Score: 15/100: The team size is estimated to be under 10, with Florian Palmade and Clara Berger listed as directors (pappers.fr), but no other key executives are visible.
- Completeness (25%) | Score: 5/100: C-suite visibility is non-existent beyond the CEO, and the tech vs. commercial balance is unknown, indicating a critical gap in leadership data.

🌐 MARKET OPPORTUNITY (0%) | Score: 85/100
The market for AI developer infrastructure, specifically RAG-as-a-Service, is experiencing explosive growth as a direct tailwind from widespread LLM adoption.

- Size & Growth (25%) | Score: 90/100: The Managed RAG-as-a-Service for developers and enterprises building proprietary AI applications. market is a multi-billion dollar opportunity, with growth directly tied to the exponential rise of enterprise AI applications.
- Timing Why Now (25%) | Score: 95/100: The 'ChatGPT moment' created a massive, immediate need for enterprises to connect proprietary data to LLMs, and the high failure rate of initial DIY RAG projects creates a perfect entry point for a managed solution.
- Competition (25%) | Score: 65/100: The market is crowded with cloud providers' native tools and well-funded startups, but Laive's focus on a specialized multi-layer retrieval pipeline could be a key differentiator.
- Expansion (25%) | Score: 90/100: The company is based in France, suggesting an initial European focus with clear potential for international expansion and moves into adjacent categories like data security and AI observability.

💡 PRODUCT INNOVATION (0%) | Score: 80/100
Laive.ai presents a strong, developer-focused product that directly addresses the core complexities of building production-grade RAG systems.

- Differentiation (25%) | Score: 85/100: Key differentiators are the advanced multi-layer retrieval pipeline, sub-500ms latency claims, and a fully managed API approach that contrasts with more DIY open-source tools (laive.ai).
- Product-Market Fit (25%) | Score: 70/100: Early signs are positive with claims of over 50 clients and 2M+ retrievals performed, suggesting the product is resonating with its target developer audience (laive.ai).
- Scalability (25%) | Score: 85/100: Built as a SaaS platform with a REST API and integrations for LangChain and major cloud storage providers, the architecture appears designed for scale (laive.ai).
- IP & Barriers (25%) | Score: 80/100: While no patents are mentioned, the GDPR-compliant, end-to-end encryption and potential for on-prem deployment create critical barriers for enterprise adoption, particularly in Europe.

💼 BUSINESS MODEL (0%) | Score: 75/100
The company employs a classic, proven PLG-to-enterprise SaaS model with clear monetization tiers.

- Unit Economics (25%) | Score: 80/100: Pricing is transparent and publicly available, with a freemium tier to drive adoption and clear upgrade paths, demonstrating a solid understanding of PLG economics (laive.ai).
- Revenue Model (25%) | Score: 75/100: The model is a standard recurring SaaS subscription, but the split between self-serve and enterprise revenue is currently unknown.
- Monetization (25%) | Score: 75/100: The four-tiered pricing structure (Free, Pro, Team, Enterprise) provides clear upsell paths based on usage (pages, tokens) and features (workspaces, support).
- Capital Efficiency (25%) | Score: 70/100: With a listed capital of just 4500 euros and a small team, the company appears capital-efficient, likely supplemented by non-dilutive funding from France 2030 (pappers.fr, LinkedIn).

📈 TRACTION & GROWTH (0%) | Score: 60/100
For a company founded in 2024, there are tangible, albeit self-reported, signs of early-market traction.

- Revenue Growth (25%) | Score: 50/100: While specific revenue is unknown, securing non-dilutive funding from France 2030 is a positive signal for a pre-seed company (LinkedIn).
- Customer Validation (25%) | Score: 65/100: The claim of 'Over 50 clients' is a strong indicator of early validation, though the logos and their significance are not disclosed (laive.ai).
- KPI Progression (25%) | Score: 65/100: The headcount is estimated at <10, which is normal for this stage, and metrics like 150K+ files processed and 2M+ retrievals show product usage (laive.ai).
- Market Penetration (25%) | Score: 60/100: With a base in France, the initial market penetration appears to be European, targeting the broad vertical of AI developers across industries.

🔍 RISK TO UNDERWRITE :
The entire investment thesis rests on the assumption that an unknown, unproven founder can execute and win in a hyper-competitive market against giants like Google and funded startups like Vectara. If Florian Palmade is not a 1-in-100 outlier in terms of grit, vision, and execution speed, the technically sound product will fail to capture meaningful market share, and we will know within 12 months if they fail to convert their early developer traction into paying enterprise contracts.
This risk is only resolvable through direct, in-person diligence with the founder to assess attributes not available in public data.

LAIVE.AI'S EXECUTIVE SUMMARY (2)

 KEY COMPETITIVE ADVANTAGES :

- Fully Managed Infrastructure: Unlike open-source frameworks (LangChain, LlamaIndex) that require significant engineering overhead, Laive offers a single API call to solve RAG, drastically reducing time-to-market for developers.
- Advanced Retrieval Pipeline: The 'multi-layer retrieval' approach promises higher accuracy than simple vector similarity search offered by many vector DBs, which is a critical selling point for enterprise-grade AI agents.
- Developer-Centric & Performant: By focusing on sub-500ms latency and providing clean docs and a freemium starting point, Laive is engineered to win the hearts and minds of developers, the key decision-makers for this type of infrastructure.
- Enterprise-Ready Security: Features like GDPR compliance, end-to-end encryption, and on-premise deployment options directly address the primary adoption blockers for large, regulated enterprises.

 MOAT : WEAK

The primary moat is switching costs, built via API integration and data gravity. Once a company's production AI application is built on Laive's API and gigabytes of proprietary data have been processed and indexed, the engineering cost and operational risk of migrating to a competitor becomes a significant deterrent.

This moat compounds as more data is ingested and more internal workflows become dependent on the reliability and specific performance characteristics of Laive's retrieval engine, creating deeper system-level lock-in.

A secondary layer of defensibility does not yet exist but could emerge over time through brand recognition for best-in-class reliability and performance, becoming the default 'trusted' RAG provider.

 ASYMMETRIC WAGER

- The Bull Case:

As building in-house RAG proves to be a perpetual, resource-draining 'science project' for 99% of companies, Laive's managed API becomes the default infrastructure layer for contextual AI, crossing an inflection point where their superior retrieval accuracy makes them the indispensable 'brain stem' for enterprise AI agents, leading to an acquisition by a major cloud or data platform.

- The Bear Case :

The company's core wager—that developers will pay for a managed RAG API—is wrong because open-source frameworks like LlamaIndex evolve to become 'good enough', and the underlying vector databases (e.g., Pinecone) commoditize the retrieval layer, leaving Laive squeezed in the middle with no unique value proposition, a reality that would become evident if their free-to-paid conversion rate stalls below 2%.

 RED FLAGS

- Universal Risks: The founder's background is a complete black box, which represents an unquantifiable execution risk; furthermore, the company is entering a hyper-competitive market with minimal stated capital against giants and heavily-funded startups.
- Thesis-Specific Mismatches: Our thesis requires backing founders with a clear track record of outlier execution or an earned secret; the current data provides zero evidence for either, making this a blind bet on an unknown quantity.

 FIRST MEETING PREP KIT

Given the massive team risk contrasted with a strong product and market, the first call is solely to determine if the founder is a hidden gem.

- The Investment Angle: The wager is that Florian Palmade is an undiscovered, product-obsessed technical founder who has built a superior retrieval engine and that we can uniquely help him win by surrounding him with the go-to-market and strategic capital expertise he currently lacks.

- Killer Questions for First Call :

- You have 50+ early clients, but you're competing against LangChain's community and Pinecone's war chest. What is the one, non-obvious developer acquisition channel you've discovered that is both scalable and has a higher ROI than anything your competitors are doing?
- Your professional history isn't public, so help me understand your capacity to win. Tell me about the single most difficult technical or commercial project you've ever owned, from start to finish, that ended in a measurable, successful outcome that others thought was impossible.
- Your 'Team' plan costs €15 per seat. Can you tell me what the cost-of-goods-sold is for your single most demanding customer on that plan, and at what specific usage threshold does that customer's marginal cost to you turn negative?

- First Meeting Go/No-Go Signal :

A GO signal is a crisp, verifiable, and detailed story of past outlier execution combined with a data-driven, almost obsessive, understanding of their customer acquisition costs and infrastructure metrics. A NO-GO is any hand-waving on past achievements or a lack of precise detail on the core business numbers.

 DATA CONFIDENCE : LOW

- The data is thinnest on the founder's background and execution capability; diligence must focus almost exclusively on vetting Florian Palmade and validating the self-reported traction metrics with customer references.
- DATA GAPS : Founder's full employment history • Verifiable revenue or ARR • Customer retention/churn metrics • Details on the 'France 2030' funding • Cap table structure

COMPANY INTELLIGENCE DOSSIER - URL EVIDENCE TRACKER

Purpose: Supporting documentation with comprehensive URL evidence for Investment Score Analysis
Company: Laive.ai
Data Completeness: 35/100
Assessment: ● INSUFFICIENT DATA FOR A FIRST LOOK (<70)
Calculation: (7 URLs found ÷ 20 URLs searched) × 100 = 35% completeness
Research Date: 2024-10-27 | Total URLs Found: 7

URL EVIDENCE BY SCORING CATEGORY

- 👤

TEAM EXCELLENCE | Found 2/4 data points

•

Founder-Market Fit: https://fr.linkedin.com/posts/florianpalmade_laiveai-est-désormais-financée-par-l-activity-7345699170856394752-179w. Used for: Identifying CEO name.

•

Track Record: . Used for: No data available.

•

Leadership: https://www.pappers.fr/entreprise/laive-933994303?utm_source=openai. Used for: Identifying company directors (Florian Palmade, Clara Berger) and legal structure.

•

Completeness: . Used for: No data available.
- 🌐

MARKET OPPORTUNITY | Found 0/4 data points

•

Size & Growth: . Used for: No specific URL provided for this data point.

•

Timing Why Now: . Used for: No specific URL provided for this data point.

•

Competition: . Used for: No specific URL provided for this data point.

•

Expansion: https://www.pappers.fr/entreprise/laive-933994303?utm_source=openai. Used for: Confirming French HQ.
- 💡

PRODUCT INNOVATION | Found 2/4 data points

•

Differentiation: <https://laive.ai/>. Used for: Analyzing product features and value proposition.

•

Product-Market Fit: <https://laive.ai/>. Used for: Sourcing claims of 50+ clients and usage metrics.

•

Scalability: https://docs.laive.ai/quick-start/quick-start/?utm_source=openai. Used for: Assessing API and integration capabilities.

•

IP & Barriers: <https://laive.ai/>. Used for: Identifying security and compliance features (GDPR).
- 💼

BUSINESS MODEL | Found 2/4 data points

•

Unit Economics: <https://laive.ai/>. Used for: Analyzing public pricing tiers.

•

Revenue Model: <https://laive.ai/>. Used for: Confirming SaaS subscription model.

•

Monetization: <https://laive.ai/>. Used for: Mapping out the tiered pricing funnel.

•

Capital Efficiency: https://www.pappers.fr/entreprise/laive-933994303?utm_source=openai. Used for: Sourcing initial company capital.
- 📈

TRACTION & GROWTH | Found 2/4 data points

•

Revenue Growth: https://fr.linkedin.com/posts/florianpalmade_laiveai-est-désormais-financée-par-l-activity-7345699170856394752-179w. Used for: Confirming non-dilutive France 2030 funding.

•

Customer Validation: <https://laive.ai/>. Used for: Sourcing the '50+ clients' claim.

•

KPI Progression: https://fr.linkedin.com/company/laive-ai?utm_source=openai. Used for: Assessing company activity and communication.

•

Market Penetration: . Used for: No specific URL provided for this data point.

WEB DATA COMPLETENESS ANALYSIS
Missing Critical URLs Based on Web Research: Founder's detailed LinkedIn profile/work history, official press releases on funding, customer case studies, third-party reviews.
URLs Successfully Found: 7 out of 20 searched
Critical Data Coverage: 35% of required data points
Research Confidence Level: LOW

LAIVE.AI'S POSITION IN THE VALUE CHAIN

In a fast memo you don't have access to the value chain analysis

VALUE PROPOSITION

Value Proposition: Building the most advanced retrieval tool for AI agents to enhance reliability and contextual intelligence through an AI-powered RAG API. Laive.ai will be the 'Stripe for RAG', providing the essential, managed infrastructure that allows any developer or company to easily connect their private data to large language models, ensuring AI assistants give accurate, context-aware answers without the company needing to build and maintain a complex data pipeline themselves. A fully managed RAG (Retrieval-Augmented Generation) API that selects precise and relevant data sources to feed Large Language Models. Provides a single, fully managed API endpoint that ingests raw enterprise data from multiple sources and uses a sophisticated multi-layer retrieval pipeline to instantly find the precise information an LLM needs to answer a query accurately.

Ideal Customer Profile (ICP): AI Developers, Enterprise Data Teams, and Companies building LLM-based applications using frameworks like LangChain. Developers at companies of all sizes are purchasing a managed API service to perform the job of reliably connecting their proprietary, unstructured data to Large Language Models.

B2B or B2C: B2B - Focused on providing infrastructure, APIs, and enterprise-grade workflows for other businesses building AI agents. B2D > SaaS.

Industry: Artificial Intelligence / SaaS Infrastructure. AI Developer Infrastructure. Developer & IT Infrastructure > AI Developer Infrastructure. Managed RAG-as-a-Service for developers and enterprises building proprietary AI applications.

Contact & Legal: Entity: Laive SAS. Address: 28 avenue des Pepinieres, 94260 Fresnes, France. Registration: RCS Paris 933 994 303. Email: florian@laive.ai. Founding: Registration number suggests 2024 (933 series). Directors: Florian Palmade, Clara Berger. Capital: 4500 euros. HQ: France, Paris area. French simplified joint-stock company.

Key Client Examples & Testimonials: Over 50 clients including 150K+ files processed and 2M+ retrievals performed.

PRODUCT FEATURES

Core Solution: A fully managed RAG (Retrieval-Augmented Generation) API that selects precise and relevant data sources to feed Large Language Models. Laive.ai provides a single, fully managed API endpoint that ingests raw enterprise data from multiple sources and uses a sophisticated multi-layer retrieval pipeline to instantly find the precise information an LLM needs to answer a query accurately. Building the most advanced retrieval tool for AI agents to enhance reliability and contextual intelligence through an AI-powered RAG API.

Feature Encyclopedia: AI-powered RAG API | Advanced Multi-layer Retrieval Pipeline | Vector Similarity Search | Intelligent Database Retrieval | Context Extraction | Sub-500ms Response Times | GDPR-compliant Storage | Encrypted Embeddings | Anonymized Metadata | Vault Management | Sync Files | On-prem Deployment.

Technical Capabilities: LangChain Integration | REST API | Multi-Cloud Support (Google Drive, OneDrive, AWS S3, Dropbox, Box, SharePoint) | Sub-500ms Latency | End-to-end Encryption | GDPR compliance | On-prem Deployment.

Use Cases: Enhancing AI agent reliability | Providing LLMs with real-time business context | Replacing complex internal RAG pipelines | Building trusted AI responses based on private enterprise data.

BUSINESS MODEL AND PRICING

Business Model Analysis: Tiered SaaS Subscription (Monthly and Yearly billing options). PLG-to-enterprise SaaS model with clear monetization tiers. Classic, proven PLG-to-enterprise SaaS model. Recurring SaaS subscription.

Revenue Streams & Pricing Tiers: Free (0 Euro) | Pro (8 Euro/mo billed yearly) | Team (15 Euro/seat/mo billed yearly, min 2 seats) | Enterprise (On demand).

Plan Features: Free: 1,000 pages, 100 tokens/mo, 1 Workspace. Pro: 1,000 pages/mo, 10,000 tokens/mo, 5 Workspaces. Team: 10,000 pages/mo, 50,000 tokens/mo, Unlimited Workspaces/Members. Enterprise: Unlimited usage and 24/7 premium support. Four-tiered pricing structure (Free, Pro, Team, Enterprise) provides clear upsell paths based on usage (pages, tokens) and features (workspaces, support).

Hidden Costs & Terms: Pro and Team plans require yearly billing for the listed prices. Team plan requires a minimum of 2 seats (30 Euro/mo min). Setup fees not mentioned.

TEAM & COMPANY CULTURE

Company Culture: Performance-driven with a focus on privacy and enterprise-grade security. Service Level Agreements (SLA) define response times for blocking anomalies.

Team Analysis: Florian (Contact point via florian@laive.ai). CEO: Florian Palmade. Directors: Florian Palmade, Clara Berger. Legal documents mention Laive SAS is a French simplified joint-stock company with 4500 euros capital. Founder-Market Fit: No data available. Track Record: No data available. Leadership: The team size is estimated to be under 10. Completeness: C-suite visibility is non-existent beyond the CEO. Founder Archetype: No data available. Pedigree Signal: No data available. Loyalty & Tenure: No data available. Commercial Fit: No data available. No work history data available. No education history data available. Grit: 0/100. Discipline: 0/100. Self-Belief: 0/100. Execution: 0/100. Leadership: 0/100. Vision and Purpose: 0/100. TOTAL FOUNDER DNA SCORE: 0/100.

Job Offers & Titles: None specifically listed in the text provided.

Estimated Headcount:

Product & Engineering: 1-5

Marketing: 1

Sales: 1

Support & IT: 1

General & Admin (G&A): 1

Total estimated <10 based on early-stage registration and capital structure.

CEO

EXECUTIVE ASSESSMENT

- Founder Archetype: No data available to determine.
- Pedigree Signal: No data available to determine.
- Loyalty & Tenure: No data available to determine.
- Commercial Fit: No data available to determine.

PROFESSIONAL NARRATIVE

No data available to construct a professional narrative.

DETAILED CAREER TIMELINE

No work history data available.

ACADEMIC BACKGROUND

No education history data available.

FOUNDER DNA & PSYCHOMETRIC ASSESSMENT

- Grit: 0/100 - No work history or personal statements are available to assess persistence through challenges or long-term commitment. Without evidence of navigating difficult market conditions or long project durations, grit cannot be evaluated.
- Discipline: 0/100 - No detailed role descriptions or performance metrics are available in the provided data. Therefore, there is no basis to evaluate accountability for outcomes, consistent effort, or a track record of compounding progress.
- Self-Belief: 0/100 - No information about career risks, non-linear moves, or shifts from secure roles into ambiguous environments is present. This makes it impossible to assess the subject's ability to operate without external validation or to take contrarian paths.
- Execution: 0/100 - The absence of specific role details, project outcomes, or examples of scaling teams/products from 0 to 1 and 1 to 10 means velocity, learning loops, and effective course-correction cannot be determined.
- Leadership: 0/100 - No job descriptions, achievements related to team building, or personal statements using "we built" language are available. Therefore, there is no evidence to assess the subject's ability to multiply the output of others.
- Vision and Purpose: 0/100 - Without a self-summary, consistent career theme, or clearly articulated mission, it is impossible to discern a deeply held "why" or the flexibility in "how" the subject approaches their career path.

TOTAL FOUNDER DNA SCORE: 0/100

•Summary Assessment: With no data provided, it is impossible to perform any meaningful assessment of this subject's potential as a founder or their psychometric profile. All scores are zero due to a complete lack of information across all categories. To make any determination, comprehensive work history, education details, and ideally, a self-summary are required. Without this, no strengths, weaknesses, or complementary co-founder needs can be identified.

LAIVE.AI'S SWOT ANALYSIS

STRENGTHS	WEAKNESSES
Over 50 clients have processed 150K+ files and 2M+ retrievals via the RAG API. Sub-500ms response times enable real-time AI agent performance. GDPR-compliant storage and end-to-end encryption meet enterprise security standards. Native LangChain integration accelerates developer adoption. France 2030 public funding provides non-dilutive R&D capital.	No founder work history or psychometric data blocks leadership evaluation. Team of under 10 people risks execution bottlenecks at scale. 4500 euro registered capital signals limited financial runway. Invitation-only beta access indicates immature product readiness. Absence of disclosed private funding exposes cash flow vulnerabilities.

OPPORTUNITIES	THREATS
Proliferating AI agents demand reliable private-data RAG solutions. Enterprises seek managed APIs to simplify internal RAG builds. French public ecosystem offers grants and AI partnerships. Multi-cloud sync features target fragmented enterprise data sources. On-prem deployment appeals to high-security regulated industries.	Pinecone and Weaviate dominate vector search with larger ecosystems. AWS Bedrock and GCP Vertex commoditize RAG infrastructure. Single founder reliance creates key-person risk. EU AI Act may raise compliance costs for French AI firms. VC funding drought hits early-stage AI infra startups.

ACTION PLAN
How to defend? Lock in GDPR and on-prem features as moats against cloud giants targeting less-regulated markets. Build proprietary multi-layer retrieval pipelines that incumbents can't match without heavy R&D. Recruit a co-founder engineer now to distribute execution risk from Florian. How to win? Double down on LangChain integrations and sub-500ms latency to capture developer mindshare in AI agent workflows. Leverage France 2030 grants to fund aggressive pricing undercuts against Pinecone. Land 5-10 enterprise logos in regulated sectors to trigger viral B2B referrals. What would be fatal? Failure to raise private capital amid 2025 VC pullback burns through France 2030 grants by Q3 2026. Florian burns out without co-founder support, stalling product iterations as competitors ship v2 retrieval tech. Enterprises defect to free AWS RAG tools when Laive misses SLA on latency. What to fix? Disclose founder pedigrees and hire sales leadership to signal scalability beyond solo operation. Raise a seed round immediately to extend runway past 4500 euro capital. Open beta fully to convert 50 clients into paying Team/Enterprise ARR.

Proplace

09/05/2026

SYNERGIES BETWEEN LAIVE.AI AND PROFUND VC

No synergies between Laive.ai and another actors have been calculated

CONVICTION FROM AN AI GENERAL PARTNER ON LAIVE.AI

Angle
This target was evaluated under the AI Workflow Automation for European Operational Teams angle — the only fund thesis angle that captured any signal from Laive.ai's product vocabulary, though even this match is partial: the company builds developer RAG infrastructure, not an ops-and-finance workflow automation product. No other angle in the fund thesis fits.

Market
The managed RAG-as-a-Service market is growing but Laive.ai is competing in the developer infrastructure sub-segment, not in the vertical SME workflow automation market the thesis targets. Pinecone, Vectara, LlamaIndex, and AWS Kendra all compete for the same developer attention. No traction among French SME operations or finance buyers is visible.

Timing
Timing for the developer RAG market is acceptable, but premature for the workflow automation vertical this fund requires. No regulatory or competitive catalyst is forcing French SME ops buyers toward AI workflow tools, and commoditization risk closes any differentiation window within 24 months.

Company
Laive.ai is a genuine four-tiered recurring SaaS product with 50+ clients and 2M+ retrievals, but it is a developer infrastructure company with no ERP or HRIS integrations, no ops-and-finance buyer traction, and no documented roadmap for vertical workflow automation. The €4,500 stated capital reflects very early stage; financial metrics are entirely unverified.

Founder
CEO Florian Palmade has zero verifiable public background — no employment history, no academic credentials, no prior outcomes. The Founder Dependency Trap anti-pattern is triggered at maximum severity for a sub-10-person team in a commoditizing infrastructure market.

Thesis-fit
Confirmation rate is 12% against the AI Workflow Automation angle — 2 confirms (both on the universal SaaS model check), 5 direct contradictions on the highest-priority thesis requirements (no ERP integrations, no ops-and-finance buyers, no workflow roadmap), and 22 data gaps. The thesis mismatch is structural, not marginal.

Verdict
● PASS (Anti-Pattern) — Laive.ai is rejected because it is a developer RAG infrastructure product rather than an AI workflow automation platform for French SME operations and finance teams, and because the Founder Dependency Trap anti-pattern is triggered at maximum severity with zero verifiable CEO background. The Devil's Advocate correctly notes that 22 data gaps and LOW confidence mean this verdict could be revisited if a founder call and buyer-split disclosure resolve the two decisive uncertainties within 30 days.